



Hanlin Li

All That's Happening behind the Scenes: Putting the Spotlight on Volunteer Moderator Labor in Reddit

Hanlin Li¹, Brent Hecht¹, Stevie Chancellor²

¹ Northwestern University

² University of Minnesota Twin Cities

lihanlin@u.northwestern.edu, bhecht@northwestern.edu, steviec@umn.edu

Abstract

Online volunteers are an uncompensated yet valuable labor force for many social platforms. For example, volunteer content moderators perform a vast amount of labor to maintain online communities. However, as social platforms like Reddit favor revenue generation and user engagement, moderators are under-supported to manage the expansion of online communities. To preserve these online communities, developers and researchers of social platforms must account for and support as much of this labor as possible. In this paper, we quantitatively characterize the publicly visible and invisible actions taken by moderators on Reddit, using a unique dataset of private moderator logs for 126 subreddits and over 900 moderators. Our analysis of this dataset reveals the heterogeneity of moderation work across both communities and moderators. Moreover, we find that analyzing only visible work – the dominant way that moderation work has been studied thus far – drastically underestimates the amount of human moderation labor on a subreddit. We discuss the implications of our results on content moderation research and social platforms.

Introduction

Online volunteers are crucial to the success of prominent commercial social platforms, such as Reddit, Twitch, and Facebook Groups. Beyond all the publicly visible labor they do generating content, volunteers also perform managerial tasks behind the scenes such as content moderation, fact-checking, and norm-setting. This work ensures the health and vibrancy of social platforms and is essential for maintaining online communities.

Despite volunteers' utmost importance to many social platforms, they are not always the group that platforms prioritize in design and development, especially for volunteer content moderators. Prominent news outlets have reported that social platforms powered by volunteer moderators such

as Reddit and Facebook prioritize revenue-generating user engagement over meeting volunteer moderators' needs (Peck, 2019; Washington Post, 2020). Moderators often feel under-appreciated, under-supported, and under-compensated by the platforms that rely on their labor (Gilbert, 2020; Matias, 2016; Postigo, 2009). This tension between moderators and the platforms they support boils over into public disagreements and disputes, e.g. "blacking out" popular communities on Reddit by making their content private and class-action lawsuits against AOL (Centivany and Glushko, 2016; Matias, 2016; Postigo, 2009).

To properly support moderators and preserve the online communities they maintain, the design and development of social platforms must be rooted in a comprehensive understanding of this labor. Existing approaches to researching content moderation at a large scale focus primarily on moderator activities that leave public visible traces, i.e. removing content and communicating with communities publicly,¹ (Chandrasekharan et al., 2019; e.g. Fan and Zhang, 2020; Jhaver et al., 2019c). However, new research shows that additional work happens behind the scenes such as managing user behavior and maintaining community settings (Gilbert, 2020; Lo, 2018). Without accounting for moderator labor as a whole, developers and researchers of social platforms risk undervaluing and driving away these volunteers and potentially undermining their platforms.

In this paper, we seek to more completely quantify and characterize moderator behaviors on Reddit. Working with Reddit moderators directly, we collected private moderator logs, called mod logs, from over 900 moderators of 126 subreddits. Private mod logs capture many more moderator actions in addition to the publicly visible ones mentioned above. As such, our dataset allows us to study the work that has fallen through the cracks of prior work and to build a